Wannan (Winnie) Yang

New York and Seattle, USA winnieyangwn96@gmail.com personal website linkedin github

EDUCATION

Ph.D. candidate in Computational Neuroscience. New York University, Buzsáki Lab — Graduating in 2025.9
Visiting Student. MIT, Tye Lab — 2018.6 – 2019.6
B.S. in Computational Neuroscience. University of Edinburgh. *GPA: 4.0 (USA equivalent)* — 2014.9 – 2018.5

RESEARCH EXPERIENCE

LLM Interpretability and Alignment (ICLR) — 2024.3 – present

Collaborator: Chen Sun, Google DeepMind

- Designed and conducted experiments to study two safety-related problems in large language models (LLMs): *deception* and *jailbreaks*.
- Implemented a battery of interpretability tools including contrastive activation steering, activation patching and sparse auto-encoders (SAEs) to understand and control LLMs.
- Built a pipeline (github) to evaluate, analyze and steer 25+ large language models from different model families (Gemma, Llama, Pythia, Qwen and Yi) of different sizes (form 1.5 billion to 70 billion parameters).
- Published a series of technical blog posts to share the research findings.

Memory Representation and Consolidation (Science, Nature, NeurIPS) — 2020.9 – present

Mentor: György Buzsáki, NYU

- Led a project to study a key mechanism for selective memory consolidation in the brain. This novel discovery has led to a publication in *Science* (leading author).
- Developed a novel latent-space based decoding method and applied various ML tools (including Bayesian decoding) to decode the content of memory reactivations ('replays') from neural population activity during learning and sleep.
- Created a pipeline for decoding large-scale ($50TB$) electrophysiology data. The pipeline can be applied to a diverse range of tasks. Implemented variants of the method to different datasets and projects, which enabled further key publications, including a collaboration project recently accepted at *Nature* (in press) and a first author paper at *NeurIPS* Symmetry and Geometry in Neural Representations Workshop.
- Open-sourced demo codes and tutorials. The neural data processing and decoding pipeline has been been widely used by lab members and colleges from other research labs.

Brain-inspired Deep Reinforcement Learning (NeurIPS) — 2021.3 – 2023.9

Collaborator: Chen Sun, Google DeepMind

- Co-developed a brain-inspired (memory consolidation and reflection) framework to build a novel deep RL algorithm.
- The resulting simple and scalable algorithm greatly improved long-term credit assignments in a diverse set of RL tasks (including grid-world, Montezuma's Revenge and other Atari games).

PUBLICATIONS

- W. Yang, C. Sun, G. Buzsáki. (2024). **Interpretability for Safe AI: Jailbreak as a case study**. *In preparation.*
- W. Yang, G. Buzsáki. (2024). **Interpretability of LLMs Deception: Universal Motif.** *ICLR* (under review).
- W. Yang, C. Sun, R. Huszár, T. Hainmueller, K. Kiselev, G. Buzsáki. (2024). **Selection of experience for memory by hippocampal sharp wave ripple.** *Science* 383, 1478-1483.
- I. Zutshi, A. Apostolelli, W. Yang, Z. Zheng, T. Dohi, E. Balzani, A. H. Williams, C. Savin, G. Buzsáki. (2024). **Hippocampal neuronal activity is aligned with action plans.** *Nature* (in press).
- C. Sun, W. Yang, T. Jiralerspong, D. Malenfant, B. Alsbury- Nealy, Y. Bengio, B. Richards. (2023). **Contrastive Retrospection: honing in on critical steps for rapid learning and generalization in RL.** *NeurIPS.*
- W. Yang, C. Sun, R. Huszár, G. Buzsáki. (2023). **Changes in the geometry of hippocampal representations across brain states.** Symmetry and Geometry in Neural Representations Workshop *NeurIPS.*
- E. Y. Kimchi , A. Burgos-Robles, G. A. Matthews, T. Chakoma, M. Patarino, J. Weddington, C. A. Siciliano, W. Yang, S. Foutch, R. Simons, M. Fong, M. Jing, Y. Li, D. B. Polley, Kay M. Tye. (2023). **Reward contingency gates selective cholinergic suppression of amygdala neurons.** eLife
- S. Tennant , I. Hawes, H. Clark, W. Tam, J. Hua, W. Yang, K. Gerlei, E. Wood, M. Nolan. (2022). **Analogue representation of a spatial memory by ramp-like neural activity in retrohippocampal cortex.** *Current Biology*
- C. Sun, W. Yang, J. Martin, S. Tonegawa. (2020). **Hippocampal neurons represent events as transferable units of experience.** *Nature Neuroscience*.

SKILLS

ML: Pytorch, scikit-learn

LLM Agent: LangChain, AutoGen

LLM Interpretability: transformer-lens, Hugging Face Transformers, Contrastive Activation Steering, Activation Patching, SAE Steering

Programming: Python, MATLAB, HTML, LaTeX

Computational Neuroscience: Large-scale High-dimensional Data Analysis, Linear and Nonlinear Dimensionality Reduction, Time Series Data Analysis, Neural Data Decoding

COURSES

Large Language Model Agents Ongoing
Instructer: Dawn Song

Deep Learning NYU. Grade: A
Instructer: Yann LeCun

Computational Cognitive Modeling NYU. Grade: A
Instructer: Brenden Lake

Reinforcement Learning UCL.
Instructor: David Silver

Neural Circuits and Computational Modeling NYU. Grade: A
Instructor: Xiaojing Wang

Neural Networks and Deep Learning deeplearning.ai
Instructor: Andrew Ng

Applied Machine Learning UoE. Grade: A
Instructor: Oisin Mac Aodha